

Image Classification and Visualizations with Convolutional Neural Networks – ImageNet10

Name	Shilpa Gopalakrishna
------	----------------------

QUESTION I


1.1 Single-batch training

1.1.1. Display graph 1.1.1 (training & validation loss over training epochs) and briefly explain what is happening and why?

Explanation:

- Below is the chart of a simple linear model's Train/Validation loss & accuracy

Graph 1.1.1



- Linear model with no hidden layers, the input image tensor (128*128*3) is passed to Linear layer and classes are predicted as output
- Batch size=48
- Trained the model on 1 batch of train data

Brief on the task:

- Here Linear model with no hidden layers is used
- The training of model is done over single train batch of 48 records

- What is happening?**

- The model is overfitting. As we can see from the chart that the Train loss reduces very fast and is almost flat, but the validation loss is too high & also the validation accuracy is too low. Hence, we can say that the model is not able to generalize and is over fitting (train accuracy too high, train loss almost 0 but the validation loss too high & validation accuracy too low)
- **Why is the model overfitting?**
 - For a problem like image classification, linear model is too simple and is not able to recognize patterns in the image because of the linear function. Also, single batch of train set for training is not enough for the model to learn the various image patterns.
 - The model training can be stopped at 14th or 15th epoch as the training loss is flat


Epoch: 0	T_Loss: 2.8423	T_Acc: 68.7500	V_Loss: 18.3225	V_Acc: 15.9356
Epoch: 1	T_Loss: 2.1669	T_Acc: 85.4167	V_Loss: 25.5245	V_Acc: 17.1016
Epoch: 2	T_Loss: 1.3379	T_Acc: 95.8333	V_Loss: 29.3966	V_Acc: 16.8240
Epoch: 3	T_Loss: 1.0000	T_Acc: 97.9167	V_Loss: 31.3202	V_Acc: 16.2132
Epoch: 4	T_Loss: 0.7635	T_Acc: 97.9167	V_Loss: 32.8066	V_Acc: 14.2699
Epoch: 5	T_Loss: 0.4280	T_Acc: 97.9167	V_Loss: 36.0638	V_Acc: 13.7146
Epoch: 6	T_Loss: 0.0682	T_Acc: 95.8333	V_Loss: 38.0536	V_Acc: 13.1038
Epoch: 7	T_Loss: 0.0216	T_Acc: 100.0000	V_Loss: 38.3361	V_Acc: 12.8262
Epoch: 8	T_Loss: 0.0008	T_Acc: 100.0000	V_Loss: 37.5193	V_Acc: 12.9373
Epoch: 9	T_Loss: 0.0021	T_Acc: 100.0000	V_Loss: 38.6971	V_Acc: 13.2149
Epoch: 10	T_Loss: 0.0059	T_Acc: 100.0000	V_Loss: 37.6200	V_Acc: 13.4370
Epoch: 11	T_Loss: 0.0060	T_Acc: 100.0000	V_Loss: 37.0957	V_Acc: 13.5480
Epoch: 12	T_Loss: 0.0031	T_Acc: 100.0000	V_Loss: 39.4311	V_Acc: 13.2704
Epoch: 13	T_Loss: 0.0013	T_Acc: 100.0000	V_Loss: 40.1329	V_Acc: 13.2704
Epoch: 14	T_Loss: 0.0005	T_Acc: 100.0000	V_Loss: 41.2002	V_Acc: 13.4925
Epoch: 15	T_Loss: 0.0003	T_Acc: 100.0000	V_Loss: 39.7169	V_Acc: 13.2704
Epoch: 16	T_Loss: 0.0002	T_Acc: 100.0000	V_Loss: 40.8549	V_Acc: 13.0483
Epoch: 17	T_Loss: 0.0001	T_Acc: 100.0000	V_Loss: 39.3309	V_Acc: 13.0483
Epoch: 18	T_Loss: 0.0002	T_Acc: 100.0000	V_Loss: 39.8164	V_Acc: 13.0483
Epoch: 19	T_Loss: 0.0005	T_Acc: 100.0000	V_Loss: 41.2632	V_Acc: 13.1594

1.1.2 Display graph 1.1.2 (training & validation loss over training epochs, with modified architecture) and explain how and why it shows that the model is overfitting the training batch.

Explanation:

- Below is the chart of model with hidden layers Train/Validation loss & accuracy

Graph 1.1.2



Brief on the task:

- Here the model architecture contains 3 hidden layers and 2 fully connected layers
- The training of model is done over single train batch of 48 records
- Why is the model overfitting?**
 - Single batch of train set for training is not enough for the model to learn the various image patterns as the model generally ends up overfitting and fails to generalize.
 - The model would have done a bit better even without optimization if we had trained the model with a more number of records because more the data used in training, the model adjusts the weights for more number & variety of records
- How to say that the model is overfitting?**
 - The train loss is reducing, and train accuracy reached 100% which is good, but as we have trained on only one batch of train data, the validation accuracy is bad.
 - As the model is doing well on training set after evaluation & is doing bad on predictions on unseen records from validation set, we can say that the model is overfitting

Epoch: 0	T_Loss: 2.2291	T_Acc: 18.7500	V_Loss: 2.3138	V_Acc: 9.9944
Epoch: 1	T_Loss: 2.2461	T_Acc: 18.7500	V_Loss: 2.5484	V_Acc: 9.9944
Epoch: 2	T_Loss: 2.1735	T_Acc: 39.5833	V_Loss: 2.3209	V_Acc: 15.2693
Epoch: 3	T_Loss: 2.1331	T_Acc: 56.2500	V_Loss: 2.2722	V_Acc: 17.7679
Epoch: 4	T_Loss: 2.0353	T_Acc: 35.4167	V_Loss: 2.2296	V_Acc: 12.2710
Epoch: 5	T_Loss: 1.9335	T_Acc: 25.0000	V_Loss: 2.2634	V_Acc: 10.2721
Epoch: 6	T_Loss: 1.8000	T_Acc: 41.6667	V_Loss: 2.2101	V_Acc: 15.2138
Epoch: 7	T_Loss: 1.6778	T_Acc: 50.0000	V_Loss: 2.1615	V_Acc: 17.9900
Epoch: 8	T_Loss: 1.5489	T_Acc: 56.2500	V_Loss: 2.1813	V_Acc: 18.2676
Epoch: 9	T_Loss: 1.4040	T_Acc: 62.5000	V_Loss: 2.2826	V_Acc: 19.7113
Epoch: 10	T_Loss: 1.2209	T_Acc: 70.8333	V_Loss: 2.2305	V_Acc: 20.7107
Epoch: 11	T_Loss: 1.0097	T_Acc: 77.0833	V_Loss: 2.2656	V_Acc: 19.8223
Epoch: 12	T_Loss: 0.8079	T_Acc: 85.4167	V_Loss: 2.5495	V_Acc: 19.3781
Epoch: 13	T_Loss: 0.6609	T_Acc: 81.2500	V_Loss: 2.5682	V_Acc: 20.0444
Epoch: 14	T_Loss: 0.6320	T_Acc: 81.2500	V_Loss: 3.0460	V_Acc: 19.3781
Epoch: 15	T_Loss: 0.4683	T_Acc: 87.5000	V_Loss: 3.4735	V_Acc: 18.5453
Epoch: 16	T_Loss: 0.5843	T_Acc: 70.8333	V_Loss: 4.1470	V_Acc: 17.4348
Epoch: 17	T_Loss: 0.4766	T_Acc: 79.1667	V_Loss: 3.6489	V_Acc: 22.0988
Epoch: 18	T_Loss: 0.6254	T_Acc: 87.5000	V_Loss: 4.0402	V_Acc: 22.0433
Epoch: 19	T_Loss: 0.4411	T_Acc: 89.5833	V_Loss: 4.2379	V_Acc: 21.2104
Epoch: 20	T_Loss: 0.1433	T_Acc: 97.9167	V_Loss: 4.3378	V_Acc: 21.3770
Epoch: 21	T_Loss: 0.1690	T_Acc: 95.8333	V_Loss: 4.9936	V_Acc: 20.1555
Epoch: 22	T_Loss: 0.2762	T_Acc: 87.5000	V_Loss: 5.4383	V_Acc: 18.4897
Epoch: 23	T_Loss: 0.1092	T_Acc: 95.8333	V_Loss: 4.5784	V_Acc: 19.9334
Epoch: 24	T_Loss: 0.0616	T_Acc: 100.0000	V_Loss: 4.1861	V_Acc: 20.6552
Epoch: 25	T_Loss: 0.0817	T_Acc: 97.9167	V_Loss: 4.2994	V_Acc: 21.0994
Epoch: 26	T_Loss: 0.1090	T_Acc: 100.0000	V_Loss: 4.3500	V_Acc: 20.8218
Epoch: 27	T_Loss: 0.0923	T_Acc: 100.0000	V_Loss: 4.3195	V_Acc: 20.8773
Epoch: 28	T_Loss: 0.0485	T_Acc: 100.0000	V_Loss: 4.4505	V_Acc: 20.9883
Epoch: 29	T_Loss: 0.0237	T_Acc: 100.0000	V_Loss: 4.4972	V_Acc: 21.2660

1.1.3 Fill in table 1.1.3 (your adjusted architecture after single-batch training), adding rows and columns as necessary.

Explanation:

Adjusted network architecture without Optimization includes:

- 3 Conv2d Layers, each with a Relu activation & a maxpool layer
- Followed by 2 Fully connected Linear layer with Relu activation
- And the final Linear layer that converts 512 features to 10 class

The filter channels extracts the features from the image and maxpool & stride reduce the image dimension at the same time making sure that the filter extracts the features from the image. At the end of the linear layer, the feature size is reduced to 512 and extracting final class from those features.

Input channels	Output channels	Layer type	Kernel size	Pad	Stride
3	32	Conv2d	3	1	1
		Relu			
		maxpool	4		2
32	40	Conv2d	2	0	1
		Relu			
		maxpool	2		2
40	56	Conv2d	3	0	2
		Relu			
		maxpool	3		2
		Flatten	7*7*56		
		Linear	7*7*56, 1024		
		Relu			
		Linear	1024, 512		
		Relu			
		Linear	512, 10		


1.2 Fine-tuning on full dataset [18 marks]

1.2.1 Display graph 1.2.1 and indicate what the optimal number of training epochs is and why.

Explanation:

- Below is the chart of finetuned model's Train/Validation loss & accuracy

Graph 1.2.1



- After finetuning the model & on training the model with full train dataset for 35 epochs, we get below accuracy & loss on 35th epoch.

- Train accuracy=82.16%, Train loss=0.5
- Validation accuracy=69%, Validation loss= 1.1

- **Evaluating the model:** We can see from the below plot that the train & validation accuracy is gradually increasing. Train loss is reducing constantly whereas validation loss is changing in minute scale.
- **Optimal number of epochs:** By increasing the epochs and training with more epochs, could have improved the test & validation accuracy. Also, we could have trained the network initially with a higher learning rate and gradually reduce the learning rate so that we could have reached the global minima sooner. With the current network, it might take >20 more epochs to reach minimal loss as the rate of change in loss is in small scale

Epoch: 0	T_Loss: 1.6810	T_Acc: 40.2639	V_Loss: 1.6366	V_Acc: 41.4214
Epoch: 1	T_Loss: 1.5107	T_Acc: 47.7361	V_Loss: 1.5124	V_Acc: 45.1971
Epoch: 2	T_Loss: 1.3809	T_Acc: 51.2639	V_Loss: 1.4676	V_Acc: 49.5836
Epoch: 3	T_Loss: 1.2580	T_Acc: 56.5139	V_Loss: 1.3040	V_Acc: 53.2482
Epoch: 4	T_Loss: 1.2855	T_Acc: 55.4306	V_Loss: 1.4134	V_Acc: 52.3043
Epoch: 5	T_Loss: 1.1128	T_Acc: 61.0000	V_Loss: 1.2154	V_Acc: 57.3570
Epoch: 6	T_Loss: 1.1178	T_Acc: 61.6528	V_Loss: 1.1921	V_Acc: 59.1338
Epoch: 7	T_Loss: 1.0203	T_Acc: 64.7778	V_Loss: 1.1353	V_Acc: 60.3554
Epoch: 8	T_Loss: 0.9964	T_Acc: 65.9861	V_Loss: 1.1632	V_Acc: 60.5775
Epoch: 9	T_Loss: 0.9772	T_Acc: 66.1528	V_Loss: 1.1269	V_Acc: 61.4659
Epoch: 10	T_Loss: 0.9227	T_Acc: 68.4167	V_Loss: 1.0338	V_Acc: 62.6319
Epoch: 11	T_Loss: 0.9603	T_Acc: 66.5972	V_Loss: 1.1717	V_Acc: 61.4103
Epoch: 12	T_Loss: 0.8675	T_Acc: 69.5000	V_Loss: 1.0608	V_Acc: 65.0750
Epoch: 13	T_Loss: 0.9403	T_Acc: 68.8472	V_Loss: 1.1249	V_Acc: 61.5214
Epoch: 14	T_Loss: 0.8389	T_Acc: 70.4583	V_Loss: 1.0780	V_Acc: 63.0205
Epoch: 15	T_Loss: 0.7601	T_Acc: 73.9167	V_Loss: 1.0406	V_Acc: 66.0189
Epoch: 16	T_Loss: 0.7525	T_Acc: 73.7639	V_Loss: 1.0001	V_Acc: 66.7407
Epoch: 17	T_Loss: 0.7451	T_Acc: 73.9583	V_Loss: 1.0772	V_Acc: 64.9639
Epoch: 18	T_Loss: 0.7966	T_Acc: 72.8194	V_Loss: 1.0831	V_Acc: 64.7418
Epoch: 19	T_Loss: 0.6996	T_Acc: 75.2083	V_Loss: 1.0528	V_Acc: 66.6297
Epoch: 20	T_Loss: 0.6389	T_Acc: 77.7222	V_Loss: 1.0224	V_Acc: 67.8512
Epoch: 21	T_Loss: 0.6816	T_Acc: 76.1111	V_Loss: 1.0158	V_Acc: 66.5741
Epoch: 22	T_Loss: 0.6370	T_Acc: 77.9028	V_Loss: 0.9547	V_Acc: 67.4625
Epoch: 23	T_Loss: 0.7804	T_Acc: 72.7083	V_Loss: 1.2269	V_Acc: 62.8540
Epoch: 24	T_Loss: 0.6720	T_Acc: 76.6528	V_Loss: 1.0481	V_Acc: 65.4636
Epoch: 25	T_Loss: 0.5488	T_Acc: 80.7361	V_Loss: 0.9755	V_Acc: 70.0722
Epoch: 26	T_Loss: 0.5403	T_Acc: 81.7639	V_Loss: 0.9295	V_Acc: 68.5175
Epoch: 27	T_Loss: 0.5556	T_Acc: 81.0972	V_Loss: 0.9843	V_Acc: 68.4064
Epoch: 28	T_Loss: 0.5847	T_Acc: 79.7639	V_Loss: 1.0268	V_Acc: 67.1849
Epoch: 29	T_Loss: 0.5208	T_Acc: 81.9167	V_Loss: 0.9965	V_Acc: 67.0738
Epoch: 30	T_Loss: 0.4836	T_Acc: 83.3611	V_Loss: 0.9486	V_Acc: 70.1277
Epoch: 31	T_Loss: 0.4551	T_Acc: 84.5972	V_Loss: 1.0133	V_Acc: 70.4609
Epoch: 32	T_Loss: 0.4454	T_Acc: 84.5556	V_Loss: 1.0686	V_Acc: 69.6835
Epoch: 33	T_Loss: 0.4627	T_Acc: 84.1250	V_Loss: 0.9814	V_Acc: 68.9617
Epoch: 34	T_Loss: 0.5162	T_Acc: 82.1667	V_Loss: 1.1698	V_Acc: 67.9067

1.2.2 Describe in detail your fine-tuning process on the complete dataset, including any adjustments you made to the network or training process to increase prediction accuracy. Explain why these adjustments increased accuracy.

Explanation:

- Model Architecture with fine-tuned network structure looks as below
 - 3 layers of Conv2D each with Relu & Maxpool layers
 - Dropout layer
 - Batch Normalization
 - Final fully connected layers

```
net_opt = nn.Sequential(  
    nn.Conv2d(3,32, kernel_size=3, stride=1, padding=1),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=4, stride=2),  
  
    nn.Conv2d(32,40, kernel_size=2, stride=1),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=2, stride=2),  
    nn.Dropout(0.05),  
  
    nn.Conv2d(40,56, kernel_size=3, stride=2),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=3, stride=2),  
    nn.Dropout(0.1),  
  
    nn.Flatten(),  
    nn.Linear(7*7*56,1024),  
    nn.BatchNorm1d(num_features=1024),  
    nn.ReLU(),  
    nn.Dropout(0.1),  
    nn.Linear(1024,512),  
  
    nn.ReLU(),  
    nn.Linear(512,10)
```

- Various optimization steps to improve performance are as below:
 1. **Added multiple layers in the network** – Added additional Conv2D layers in the network. Optimized the network by initially starting with higher kernel size with padding and gradually reducing the kernel size. Having a deeper network allows the model to learn more complex features and patterns and with optimal tuning, will provide better predictions than a shallow network.
 2. **Augmentation** – Data augmentation was done on training data by randomly flipping the training data, adding jitter & random affine. But this did not increase the accuracy much on validation data as our validation set does not include any images that are drastically different from training set in terms of angles and orientation in the image. However, this helped the model training with more variety of data that were scaled, flipped, or rotated.

```
data_transform = transforms.Compose([
    transforms.Resize(128),
    transforms.CenterCrop(128),
    transforms.ToTensor(),
    transforms.ColorJitter(hue=0.2, saturation=0.2, brightness=0.2),
    transforms.RandomAffine(degrees=10, translate=(0.1,0.1), scale=(0.9,1.1)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.Normalize(NORM_MEAN, NORM_STD),
])
```

3. **Added Dropout layers** -After fine tuning the model with additional hidden layers, dropout was introduced at 3 layers. Initially, trained the network with 2 dropout with a probability of 0.3 and 0.5, but the accuracy was too low. Hence, reduced the dropout probability to 0.1 in 2 places and introduced another dropout with 0.05 probability.
4. **Batch normalization**: Included batch normalization in the fully connected layer that normalizes the output of a layer before being fed to next layer

1.2.3 Display two confusion matrices 1.2.3 (one each for complete validation set and complete training set) for your final trained model and interpret what is shown.

Explanation:

- The train accuracy & loss on complete training data is 82.6% & 0.5 respectively
- The validation accuracy & loss on complete validation data is 68.18% & 1.02 respectively

```
loss_fulltrain, acc_fulltrain, all_pred_train, all_actual_train = evaluate_loader(train_loader, netopt_fulltrain)
loss_val, acc_val, all_pred_val, all_actual_val = evaluate_loader(valid_loader, netopt_fulltrain)
print('T_Loss: %.4f | T_Acc: %.4f | V_Loss: %.4f | V_Acc: %.4f' % (loss_fulltrain, acc_fulltrain, loss_val, acc_val))
```

T_Loss: 0.5031 | T_Acc: 82.6111 | V_Loss: 1.0261 | V_Acc: 68.1843


Confusion matrix:

- Confusion matrix can be used to pictographically represent the performance of classification model.
- The confusion matrix compares the actual label/class/target with the predicted label/class/target
- In the below confusion matrix plot, there are 10 targets where column indicates the actual class & rows indicates the predicted class

Confusion matrix for Train dataset:

- Below plot represents confusion matrix after the trained model is evaluated on training dataset


- The diagonal elements represent the number of correct predictions for a particular class/target (represented by row/column header)
- An example to understand below matrix is:
 - The class “baboon” is correctly predicted as “baboon” by the model in 591 instances ie 1st row, 1st column
 - The class “baboon” is wrongly predicted to other class 65 times, ie rows from 2nd to 10th in 1st column
 - Other class were wrong classified as “baboon” in 132 instances ie First row columns from 2nd to 10th



Confusion matrix for validation dataset:

- Below plot represents confusion matrix after the trained model is evaluated on validation dataset
- Like training set above, one example using validation test on understanding below matrix is:
 - The class “baboon” is correctly predicted as “baboon” by the model 118 times ie 1st row, 1st column
 - The class “baboon” is wrongly predicted to other class 14 times, ie rows from 2nd to 10th in 1st column
 - Other class were wrong classified as “baboon” in 77 instances ie First row columns from 2nd to 10th

```
plot_confusion_matrix_validation_data(all_pred_val,all_actual_val)
```



1.3 Evaluation and code

1.3.1 Please include `[my_student_username]_test_preds.csv` with your final submission. [8 marks]

Done

1.3.2 Please submit all relevant code you wrote for Question I in Python file `[my_student_username]_q1.py`. No need to include the config or ImageNet10 files. [13 marks]

Done

No response needed here.

QUESTION II

2.1 Preparing the pre-trained network

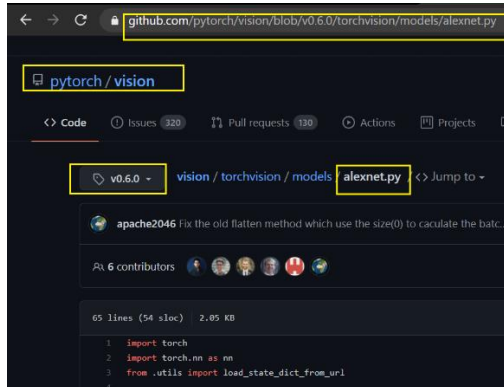
2.1.1 Read through the provided template code for the AlexNet model `alexnet.py`. What exactly is being loaded in line 59? [2 marks]

Explanation:

```
model = torch.hub.load('pytorch/vision:v0.6.0', 'alexnet', pretrained=True)
```

In the above line of code, `torch.hub.load` imports/loads a pre-trained model with pretrained weights from specific github repository.

- 1st argument - name of the GIT repository ie `pytorch/vision` with a version tag of `0.6.0`
- 2nd argument – model to be imported from the specified model repository ie “alexnet”
- 3rd argument – Load the pretrained weights from the repository



2.1.2 Write the code in *explore.py* after line 50 to read in the image specified in the variable `args.image_path` and pass it through a single forward pass of the pre-trained AlexNet model.

Code added

2.1.3 Fill in function `extract_filter()` after line 84 extracting the filters from a given layer of the pre-trained AlexNet.

Code added

2.1.4 Fill in function `extract_feature_maps()` after line 105 extracting the feature maps from the convolutional layers of the pre-trained AlexNet.

Code added

Please submit all your Question II code in a Python file `[my_student_username]_explore.py`.

No response needed here.

2.1.5 Describe in words, not code, how you ensure that your filters and feature maps are pairs; that the feature maps you extract correspond to the given filter.

Explanation:

filter: The filter index and its dimension are printed by looping the convD layers of the Alexnet model

Index 0 - torch.Size([64, 3, 11, 11]) (out_channel, in_channels, image dimension of 11*11)

Index 3 - torch.Size([192, 64, 5, 5])

Index 6 - torch.Size([384, 192, 3, 3])

Index 8 - torch.Size([256, 384, 3, 3])

Index 10 - torch.Size([256, 256, 3, 3])

feature: The feature index at same index as ConvD and its dimension are printed by looping the image after passing through a layer at certain index of the Alexnet model

Index 0 - torch.Size([1, 64, 93, 124]) (number of images, out_channels after passing through a layer, image dimension of 93*124)

Index 3 - torch.Size([1, 192, 46, 61])

Index 6 - torch.Size([1, 384, 22, 30])

Index 8 - torch.Size([1, 256, 22, 30])







Index 10 - torch.Size([1, 256, 22, 30])

- The size of the featuremap at each layer is equal to the out channels of the corresponding filter layer.
- At index 0, filter has 64 channels as output channels and the corresponding featuremap at 0th index has 64 as the number of out channels. This infers that the corresponding filter & feature maps are pairs.

2.2 Visualizations

2.2.1 For three input images of different classes, show three pairs of filters and corresponding feature maps, each from a different layer in AlexNet. Indicate which layers you chose. For each pair, briefly explain what the filter is doing (for example: horizontal edge detection) which should be confirmed by the corresponding feature map. [15 marks]

Image #1, class: Power drill

	Filter	Feature map	Brief explanation
Early layer	CONV_0 <code>torch.Size([64, 3, 11, 11])</code> 1 of 64 filters of Channel0 	<code>img shape: torch.Size([1, 3, 300, 348])</code> <code>torch.Size([1, 64, 74, 86])</code> 1 of 64 CONV2 Feature  1 of 64 CONV2 Feature after Relu 	<u>Filter layer:</u> <ul style="list-style-type: none"> • Input channels=3 • Output channels=64 • To each input channels, 64 filter is applied • Filter is identifying the vertical lines <u>Feature Map:</u> <ul style="list-style-type: none"> • By applying 64 filters on each of the 3 input channel, we get 64 feature maps (image must pass through previous layers before reaching to this layer)
Intermediate layer Layer 3	CONV_6 <code>torch.Size([384, 192, 3, 3])</code> 1 of 384 filters of Channel0 	<code>img shape: torch.Size([1, 3, 300, 348])</code> <code>torch.Size([1, 384, 17, 20])</code> 1 of 384 CONV2 Feature  1 of 384 CONV2 Feature after Relu 	<u>Filter layer:</u> <ul style="list-style-type: none"> • Input channels=192 • Output channels=384 • To each input channels, 384 filter is applied <u>Feature Map:</u> <ul style="list-style-type: none"> • By applying 384 filters on each of the 192 input channel, we get 384 feature maps (image must pass through previous layers before reaching to this layer)
Deep layer	CONV_10 <code>torch.Size([256, 256, 3, 3])</code>	<code>img shape: torch.Size([1, 3, 300, 348])</code> <code>torch.Size([1, 256, 17, 20])</code>	<u>Filter layer:</u> <ul style="list-style-type: none"> • Input channels=256 • Output channels=256


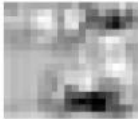

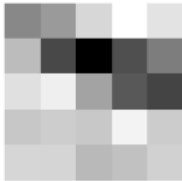
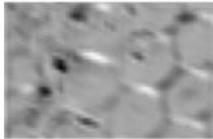
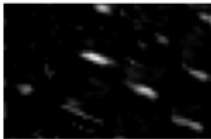
	1 of 256 filters of Channel0 	1 of 256 CONV2 Feature  1 of 256 CONV2 Feature after Relu 	<ul style="list-style-type: none"> To each input channels, 256 filter is applied <p>Feature Map:</p> <ul style="list-style-type: none"> By applying 256 filters on each of the 256 input channel, we get 256 feature maps (image must pass through previous layers before reaching to this layer)
--	---	---	--

Image #2, class: **Orange** (imagenet10/Alexnet/jybcucurxm.JPEG)

	Filter	Feature map	Brief explanation
Early layer	CONV_3 torch.Size([192, 64, 5, 5]) 1 of 192 filters of Channel0 	 1 of 192 CONV2 Feature after Relu 	<p>Filter layer:</p> <ul style="list-style-type: none"> Input channels=64 Output channels=192 To each input channels, 192 filter is applied <p>Feature Map: By applying 192 filters on each of the 64 input channel, we get 192 feature maps (image must pass through previous layers before reaching to this layer)</p>
Intermediate layer	CONV_8 torch.Size([256, 384, 3, 3]) 1 of 256 filters of Channel0	1 of 256 CONV2 Feature	<p>Filter layer:</p> <ul style="list-style-type: none"> Input channels=384 Output channels=256

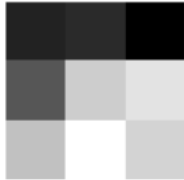
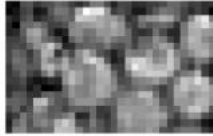
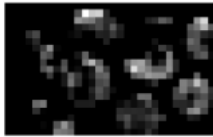

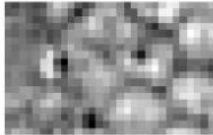
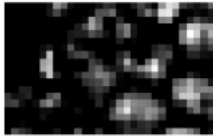
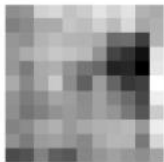

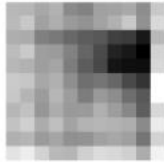
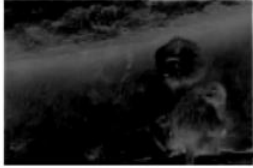
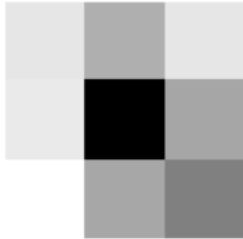
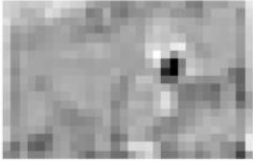
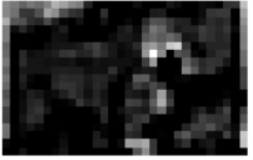

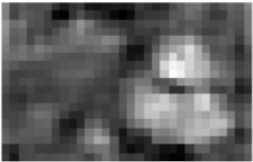
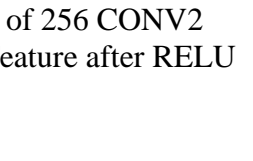
		 1 of 256 CONV2 Feature after Relu 	<ul style="list-style-type: none"> To each input channels, 256 filter is applied <p>Feature Map: By applying 256 filters on each of the 384 input channel, we get 256 feature maps (image must pass through previous layers before reaching to this layer)</p>
Deep layer	CONV_10 torch.Size([256, 256, 3, 3]) 	1 of 256 CONV2 Feature  1 of 256 CONV2 Feature after Relu 	<p>Filter layer:</p> <ul style="list-style-type: none"> Input channels=256 Output channels=256 To each input channels, 256 filter is applied <p>Feature Map: By applying 256 filters on each of the 256 input channel, we get 256 feature maps (image must pass through previous layers before reaching to this layer)</p>

Image #3, class: Baboon

	Filter	Feature map	Brief explanation
Early layer	CONV_0 torch.Size([64, 3, 11, 11]) 1 of 64 filters of Channel0  1 of 64 filters of Channel1	img shape: torch.Size([1, 3, 333, 500]) torch.Size([1, 64, 82, 124]) 1 of 64 CONV2 Feature 	<p>Filter layer:</p> <ul style="list-style-type: none"> Input channels=3 Output channels=64 To each input channels, 64 filter is applied Filter is identifying low intensity dark region surrounded

		1 of 64 CONV2 Feature after Relu 	by high intensity region <u>Feature Map:</u> By applying 64 filters on each of the 3 input channel, we get 64 feature maps
Intermediate layer	CONV_6 <code>torch.Size([384, 192, 3, 3])</code> 1 of 192 filters of Channel0 	 1 of 384 CONV2 Feature  1 of 384 CONV2 Feature after Relu	<u>Filter layer:</u> <ul style="list-style-type: none"> • Input channels=192 • Output channels=384 • To each input channels, 384 filter is applied <u>Feature Map:</u> By applying 384 filters on each of the 192 input channels of the image, we get 384 feature maps. (image must pass through previous layers before reaching to this layer)
Deep layer	CONV_10 <code>torch.Size([256, 256, 3, 3])</code> 1 of 256 filters of Channel0 	 1 of 256 CONV2 Feature  1 of 256 CONV2 Feature after RELU	<u>Filter layer:</u> <ul style="list-style-type: none"> • Input channels=256 • Output channels=256 • To each input channels, 256 filter is applied <u>Feature Map:</u> By applying 256 filters on each of the 256 input channel of the image, we get 256 feature maps. (image must pass through

			previous layers before reaching to this layer),
--	--	--	---

2.2.2 Comment on how the filters and feature maps change with depth into the network. [5 marks]

- Deeper networks recognize complex features
- Initial layers detect edges, lines & corners. Using these base features, further layers can detect shapes and in later layers the shapes are used to detect complex features.
- In our filter & feature map example, we see that the 1st convolutional layer has filters that looked like lines or stripes in different directions. When an image is passed through this filter, the feature map generated has the edges & corners highlighted.
- Initial layer feature map has a higher image dimension and a smaller number of channels. In further layers, we increase the number of features(channels) and reduce the image dimension. When the image dimension is reduced, the prominent features of the image are retained as features by increasing the number of out channels.

Marks reserved for overall quality of report. [5 marks]

No response needed here.